



TACHYON

A Reliable Memory-Centric Distributed Storage System

Haoyuan Li

October 16 @ Strata & Hadoop World NYC

Website: tachyon-project.org

Meetup: www.meetup.com/Tachyon

Outline

- Overview
 - Feature 1: Memory Centric Storage Architecture
 - Feature 2: Lineage in Storage
- Open Source
- Roadmap

Outline

- **Overview**
 - Feature 1: Memory Centric Storage Architecture
 - Feature 2: Lineage in Storage
- Open Source
- Roadmap



- Design next generation data analytics stack:
Berkeley Data Analytics Stack (BDAS)



a cluster manager making it easy to write and deploy distributed applications.



a parallel computing system supporting general and efficient in-memory execution.

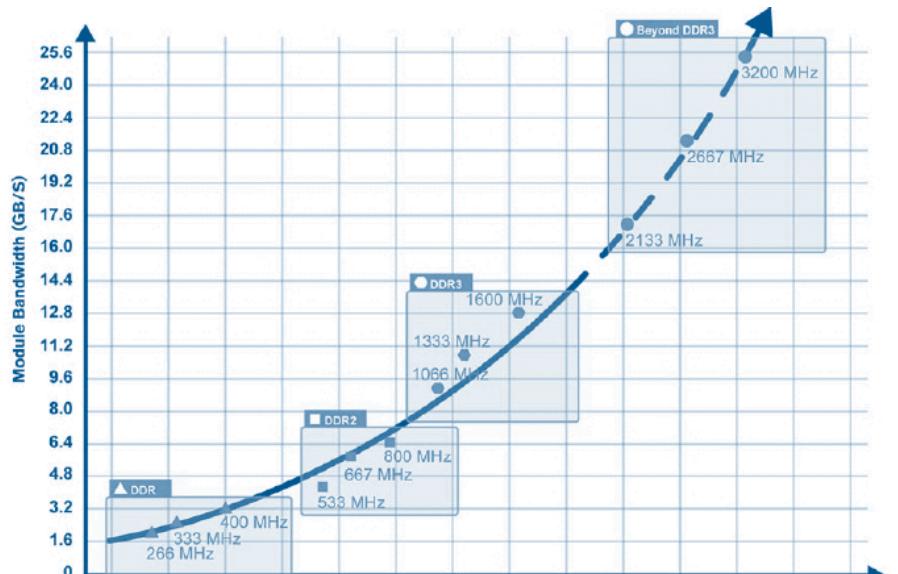


TACHYON a reliable distributed memory-centric storage enabling memory-speed data sharing.

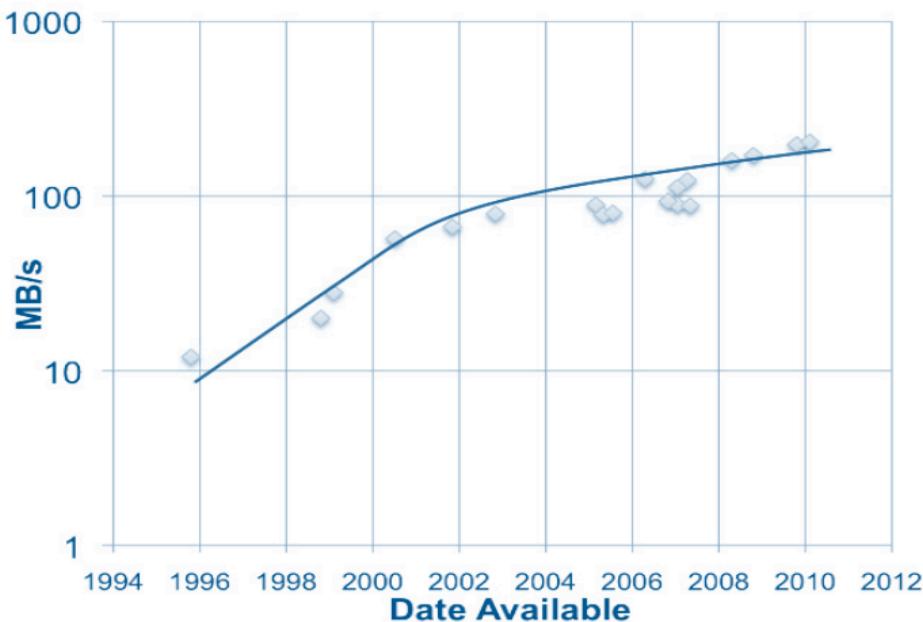
Why Tachyon?

Memory is King

- RAM throughput increasing exponentially
- Disk throughput increasing slowly



Bandwidths shown for 64-bit memory module. Date indicates approximate industry product introduction.



Memory-locality key to interactive response time

Realized by many...

- Frameworks already leverage memory

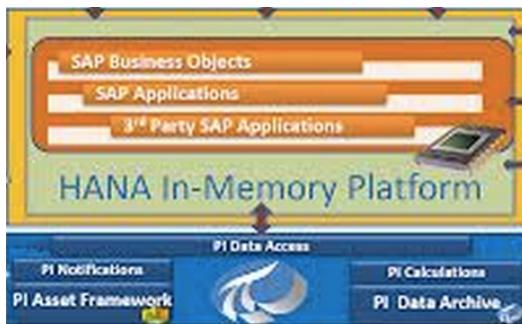


April 7, 2012

Many kinds of memory-centric data management

I'm frequently asked to generalize in some way about in-memory or memory-centric data management. I can start:

- The desire for [human real-time interactive response](#) naturally leads to

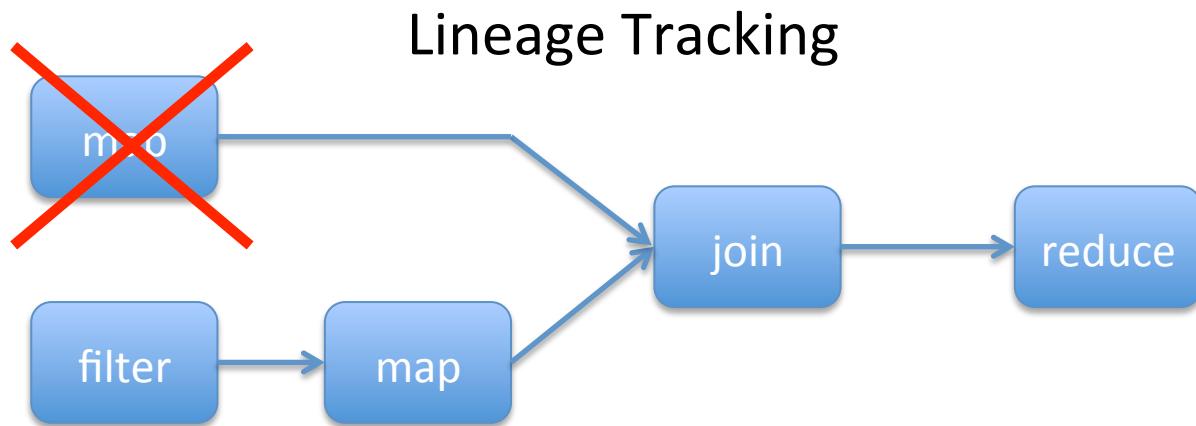


Problem solved?

Missing a Solution for Storage Layer

An Example: Spark

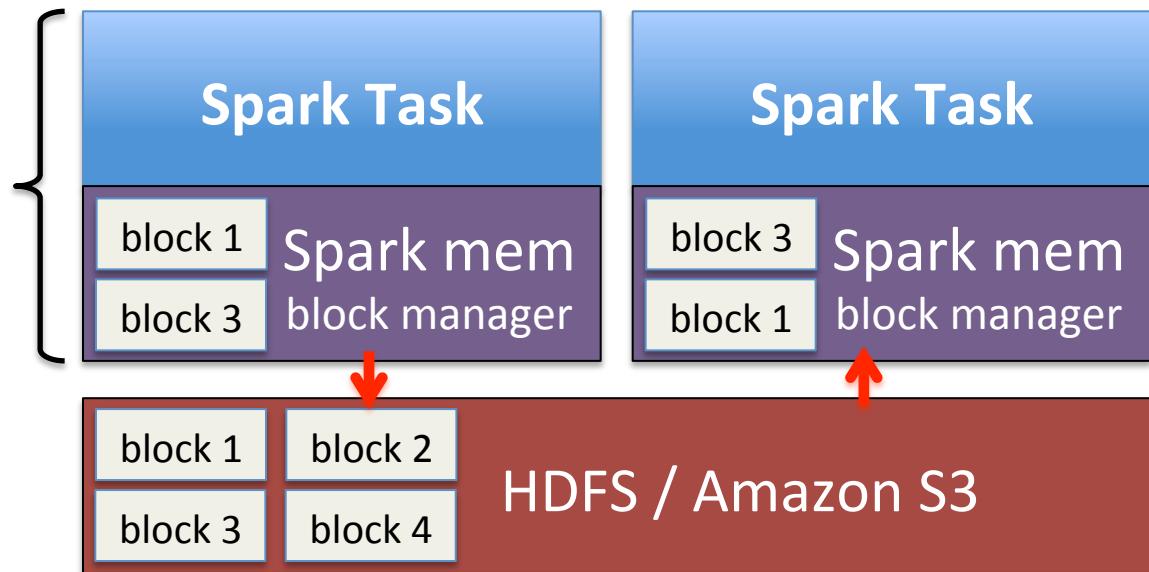
- Fast in-memory data processing framework
 - Keep **one** in-memory copy inside JVM
 - Track **lineage** of operations used to derive data
 - Upon failure, use lineage to recompute data



Issue 1

*Data Sharing is the bottleneck in analytics pipeline:
Slow writes to disk*

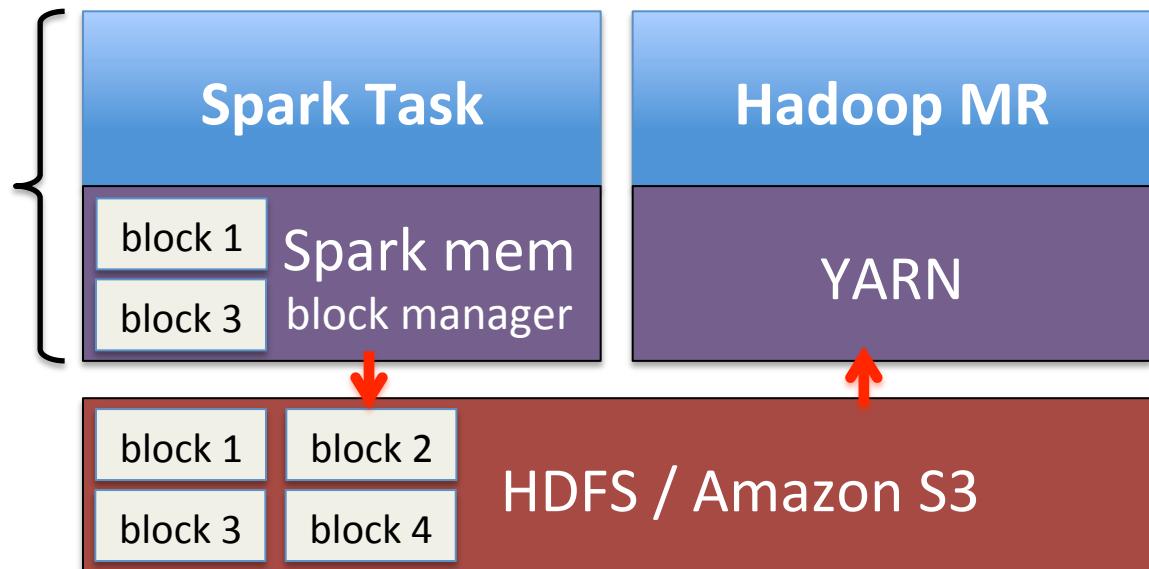
storage engine &
execution engine
same process
(slow writes)



Issue 1

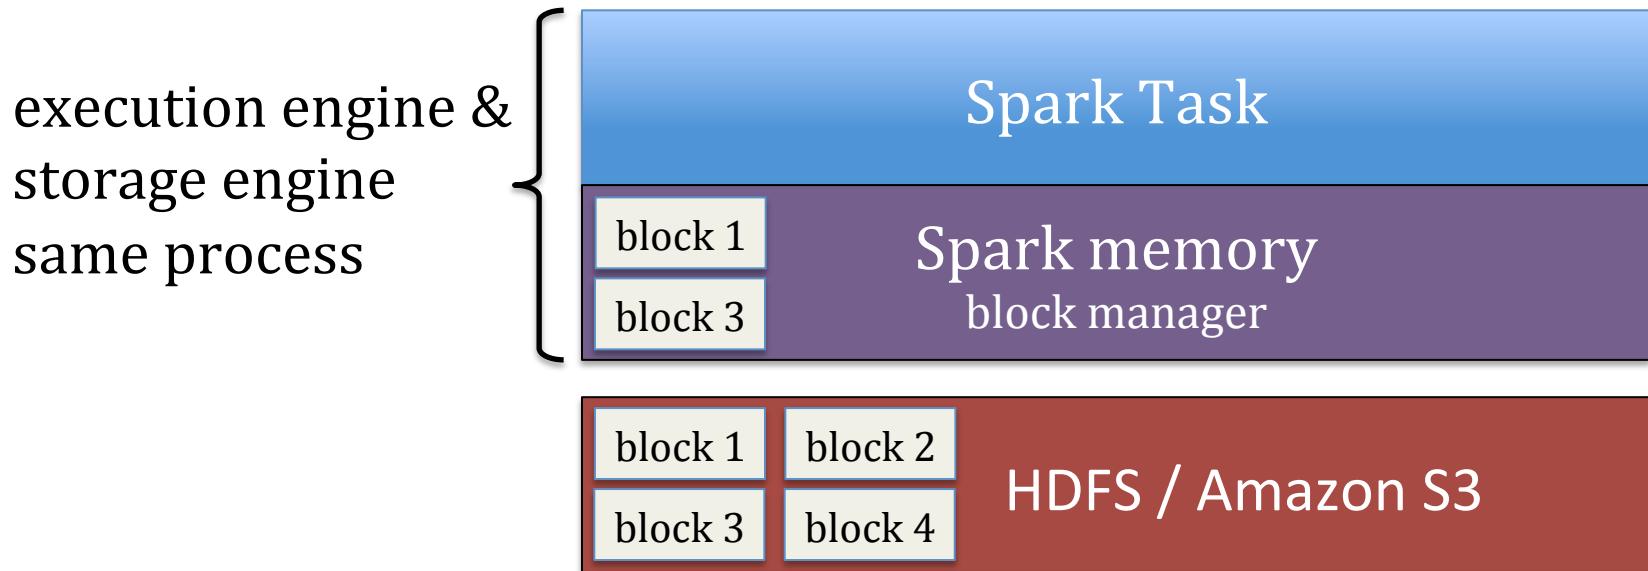
*Data Sharing is the bottleneck in analytics pipeline:
Slow writes to disk*

storage engine &
execution engine
same process
(slow writes)



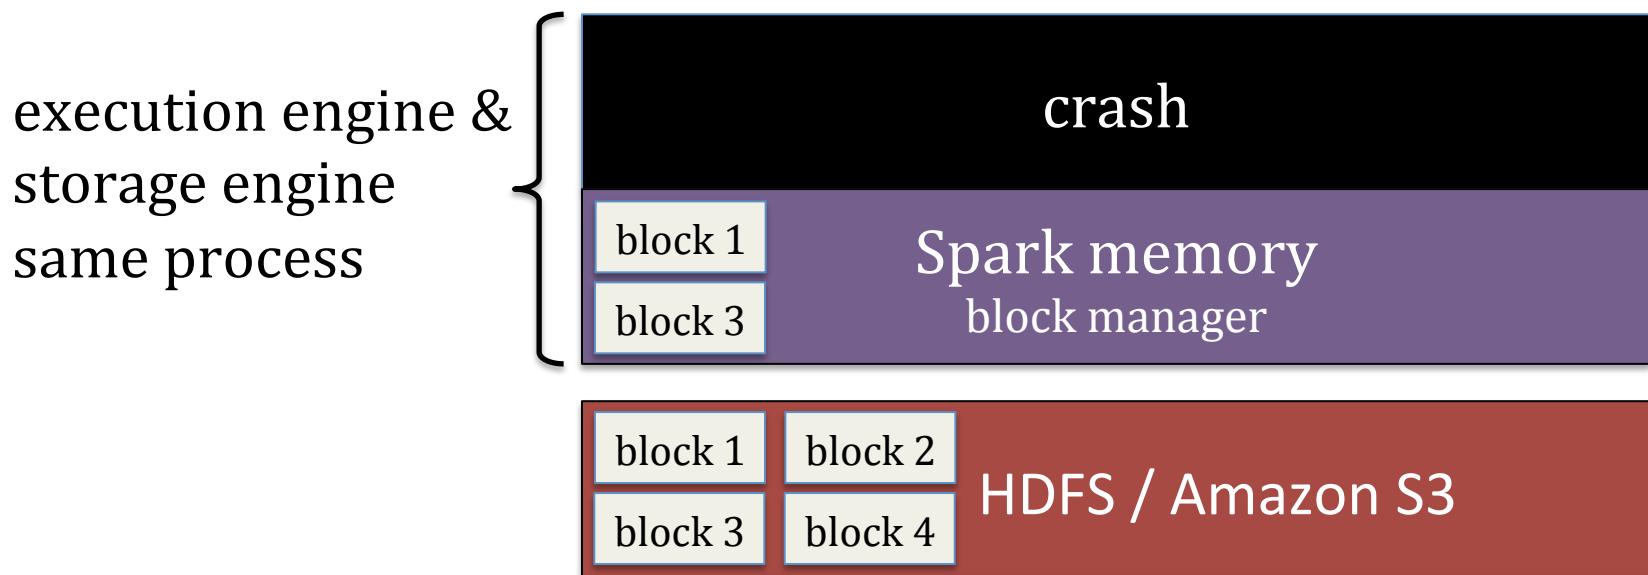
Issue 2

Cache loss when process crashes.



Issue 2

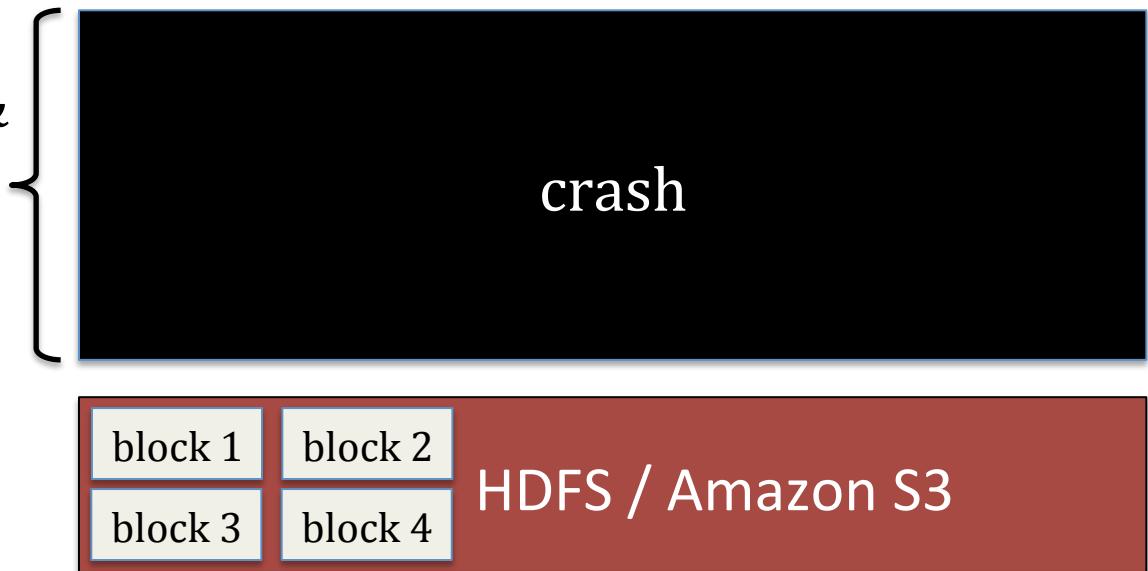
Cache loss when process crashes.



Issue 2

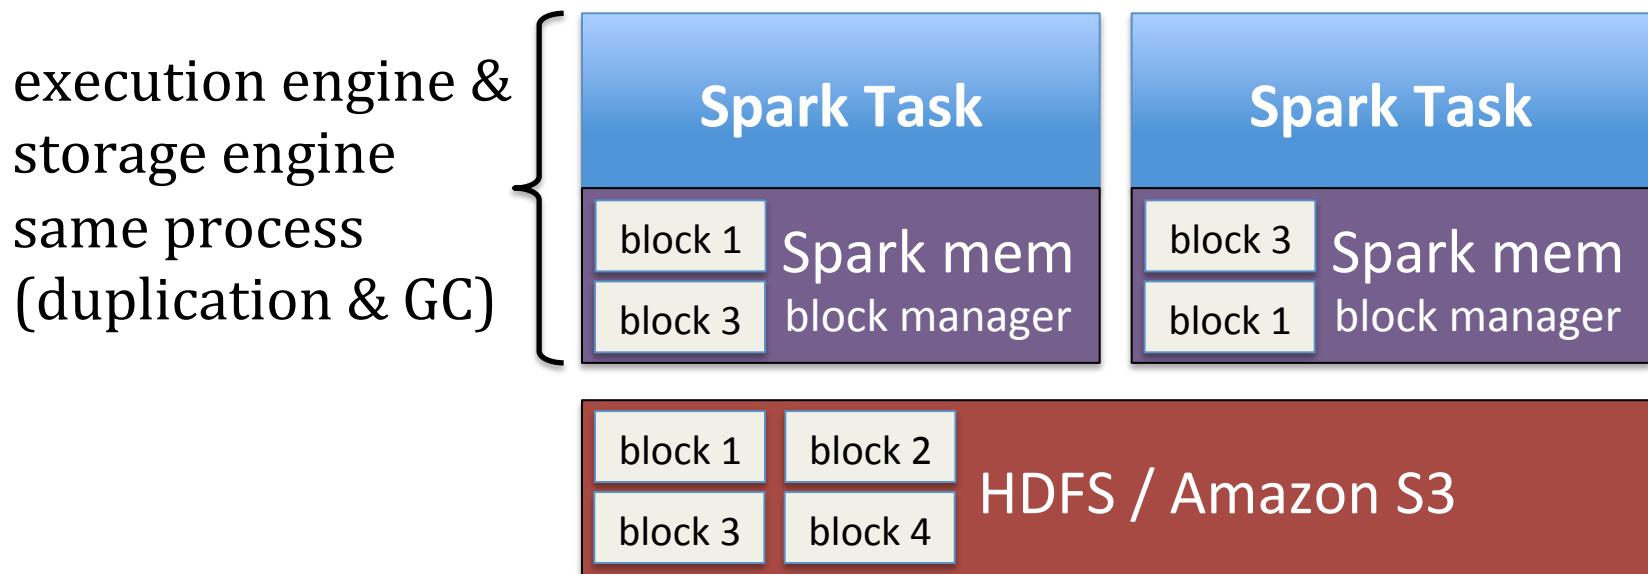
Cache loss when process crashes.

execution engine &
storage engine
same process



Issue 3

In-memory Data Duplication & Java Garbage Collection



Tachyon

Reliable data sharing at ***memory-speed***
within and across cluster frameworks/jobs

Solution Overview

Basic idea

- Feature 1: **memory-centric** storage architecture
- Feature 2: push **lineage** down to storage layer

Facts

- One data copy in memory
- Recomputation for fault-tolerance

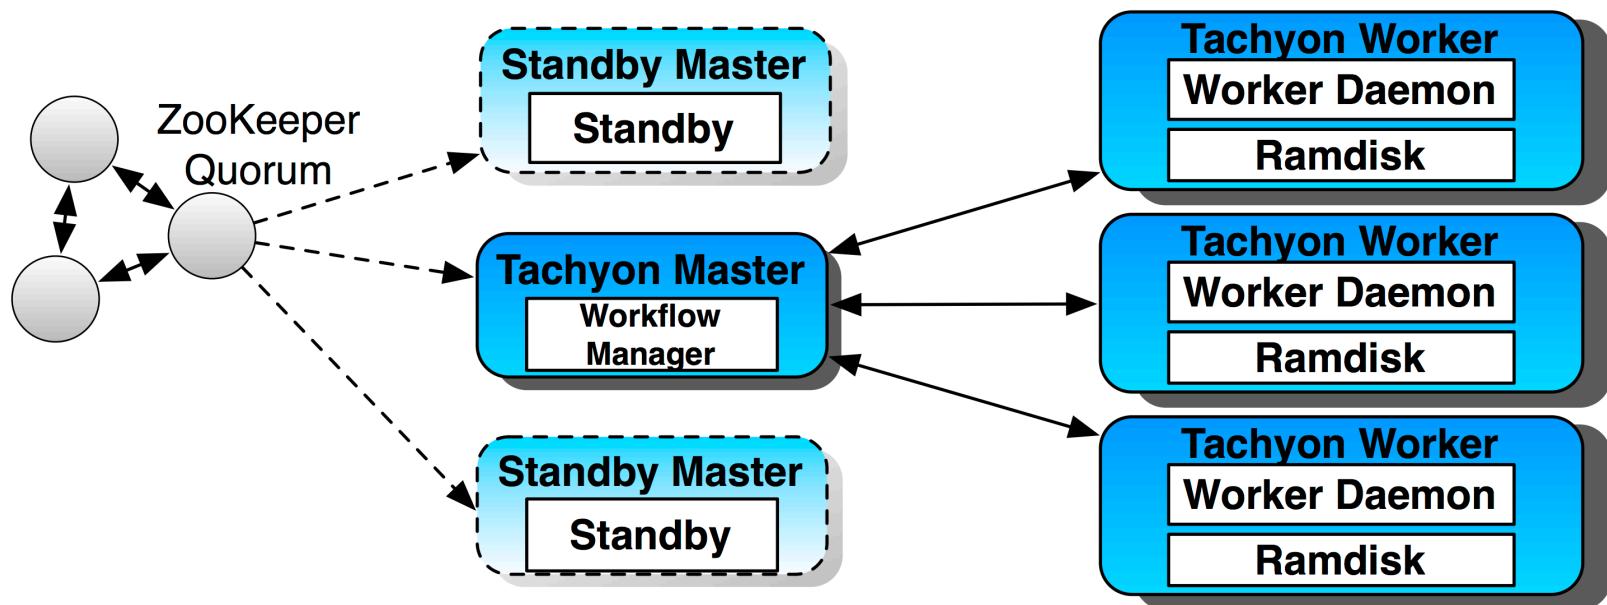
Stack

Computation Frameworks
(Spark, MapReduce, Impala, H2O, ...)

Tachyon

Existing Storage Systems
(HDFS, S3, GlusterFS, ...)

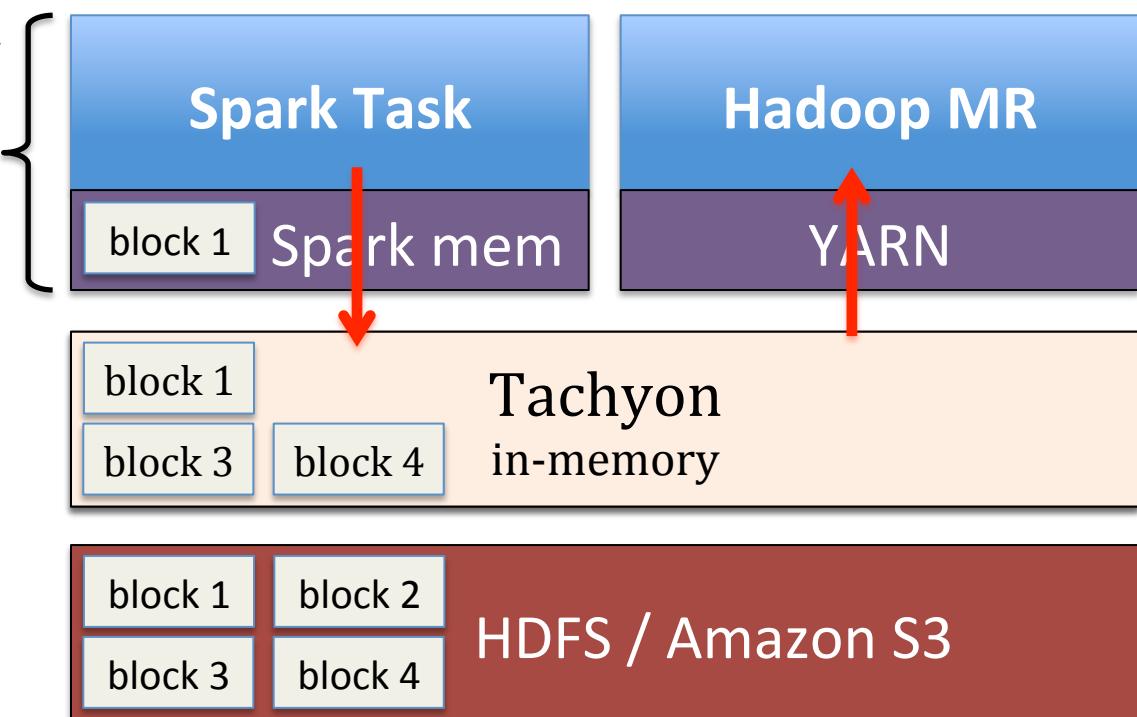
Memory-Centric Storage Architecture



Issue 1 revisited

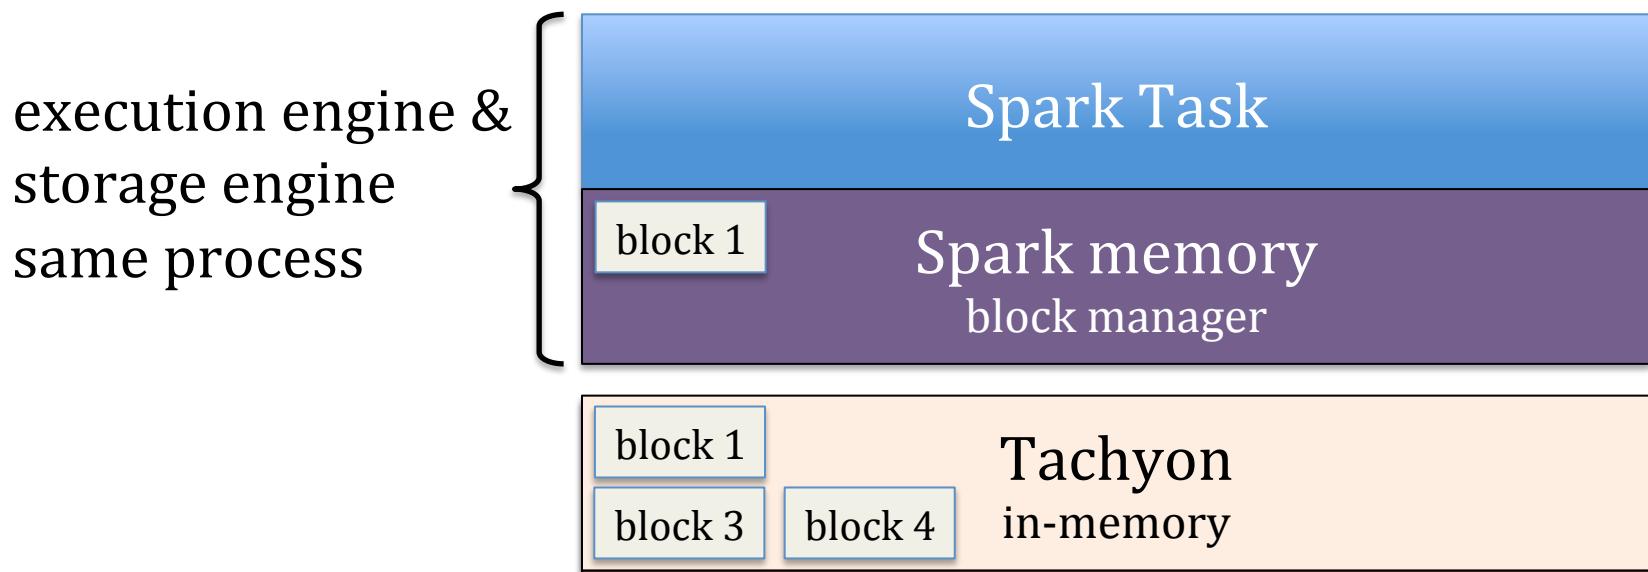
*Memory-speed data sharing
among jobs in different frameworks*

execution engine &
storage engine
same process
(fast writes)



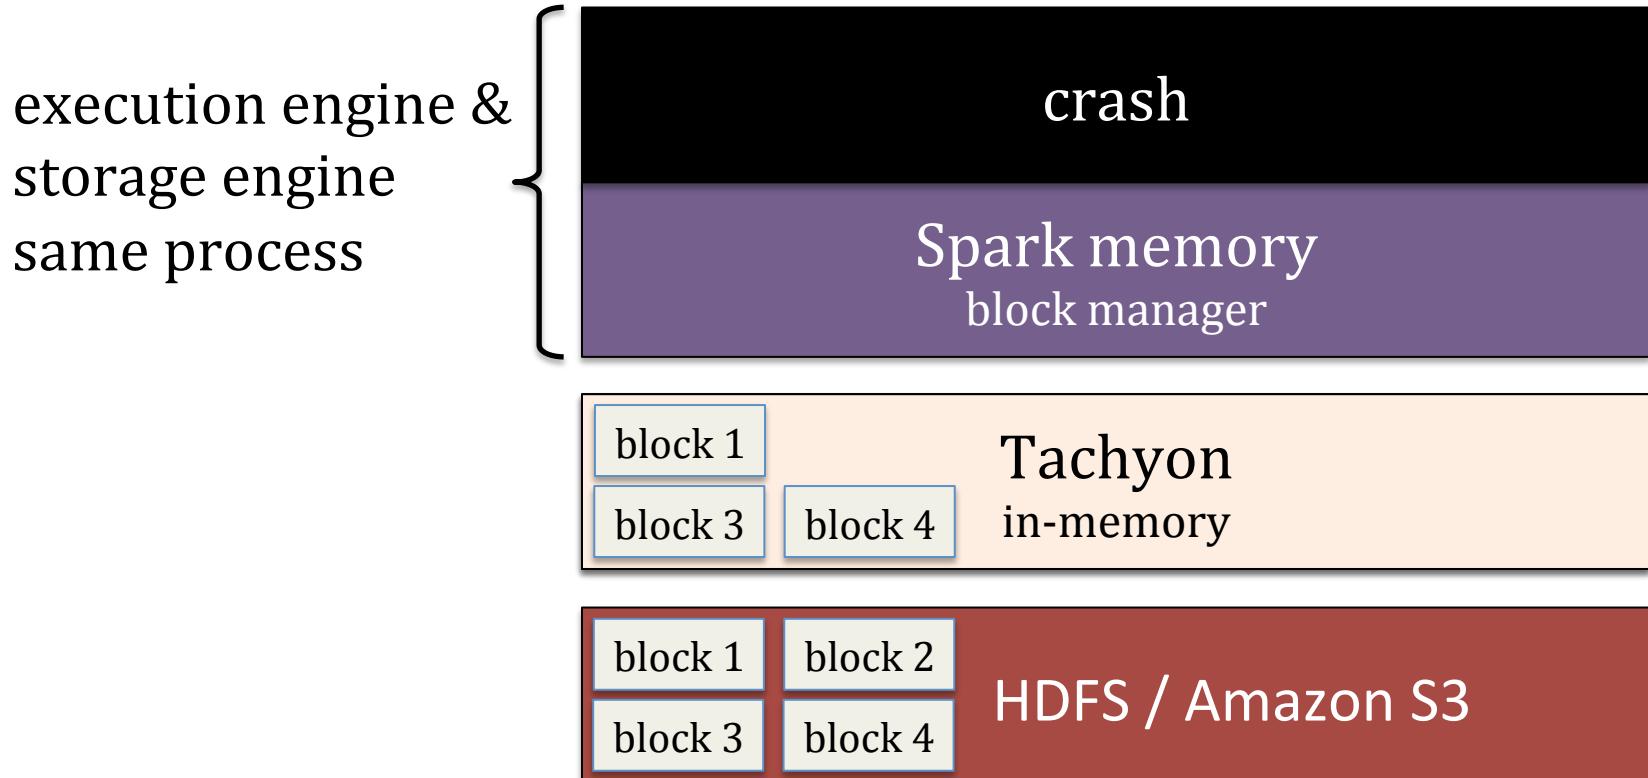
Issue 2 revisited

*Keep in-memory data safe,
even when a job crashes.*



Issue 2 revisited

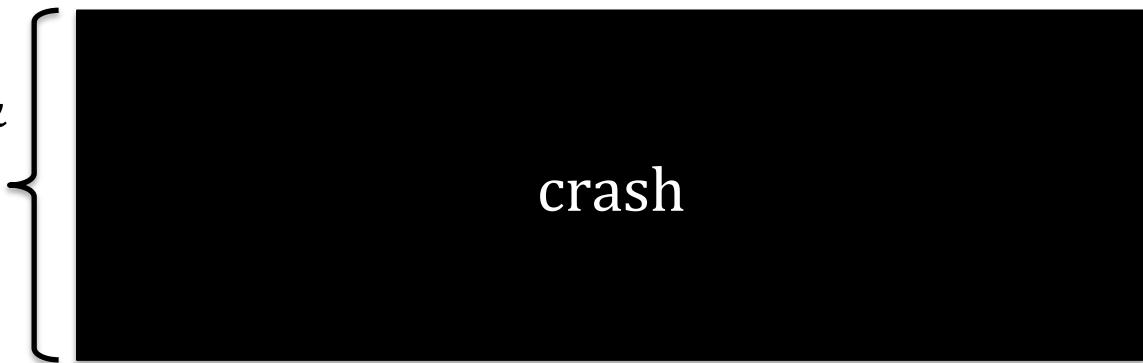
*Keep in-memory data safe,
even when a job crashes.*



Issue 2 revisited

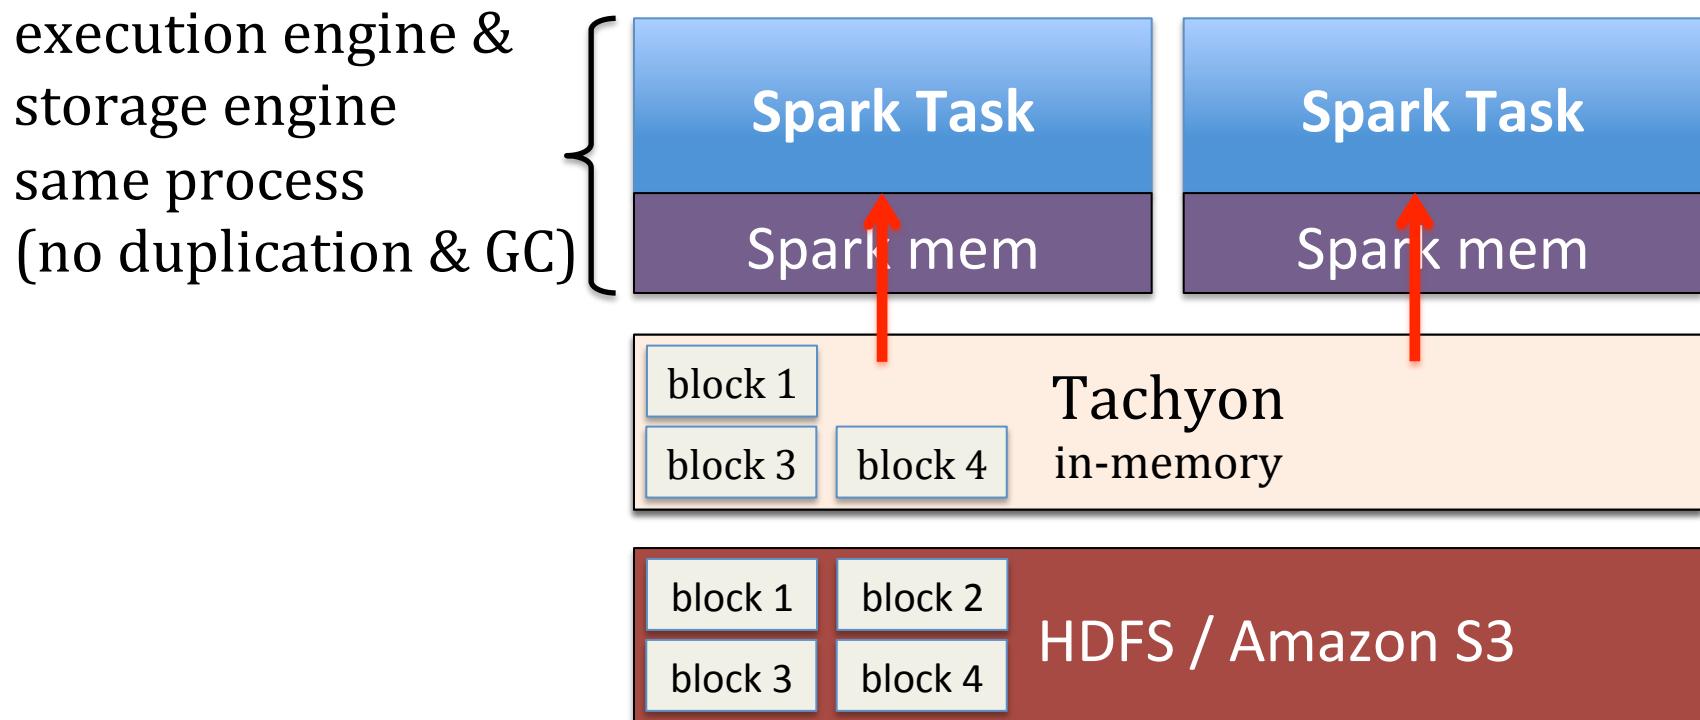
*Keep in-memory data safe,
even when a job crashes.*

execution engine &
storage engine
same process

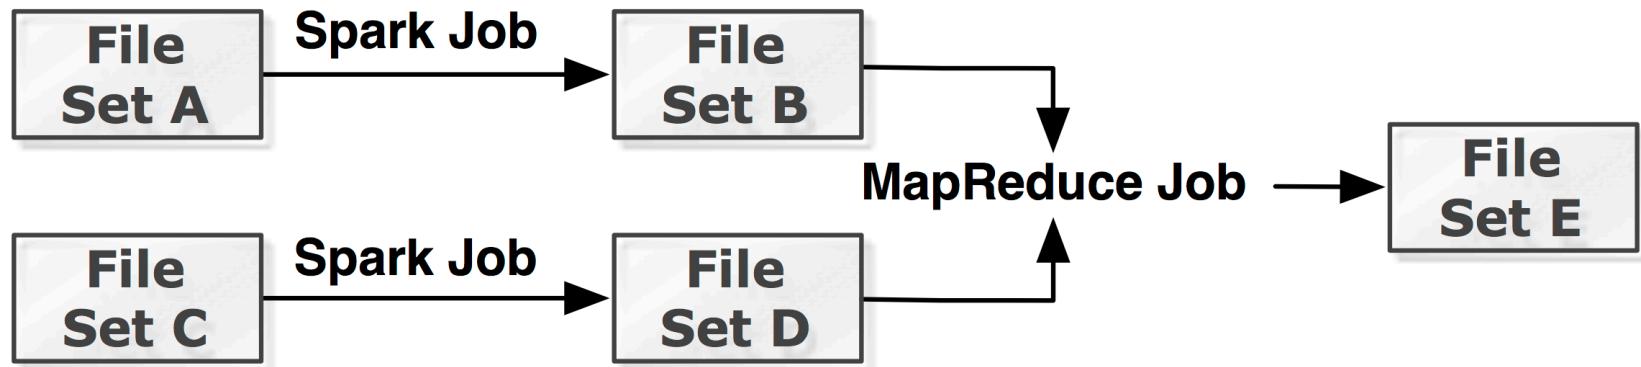


Issue 3 revisited

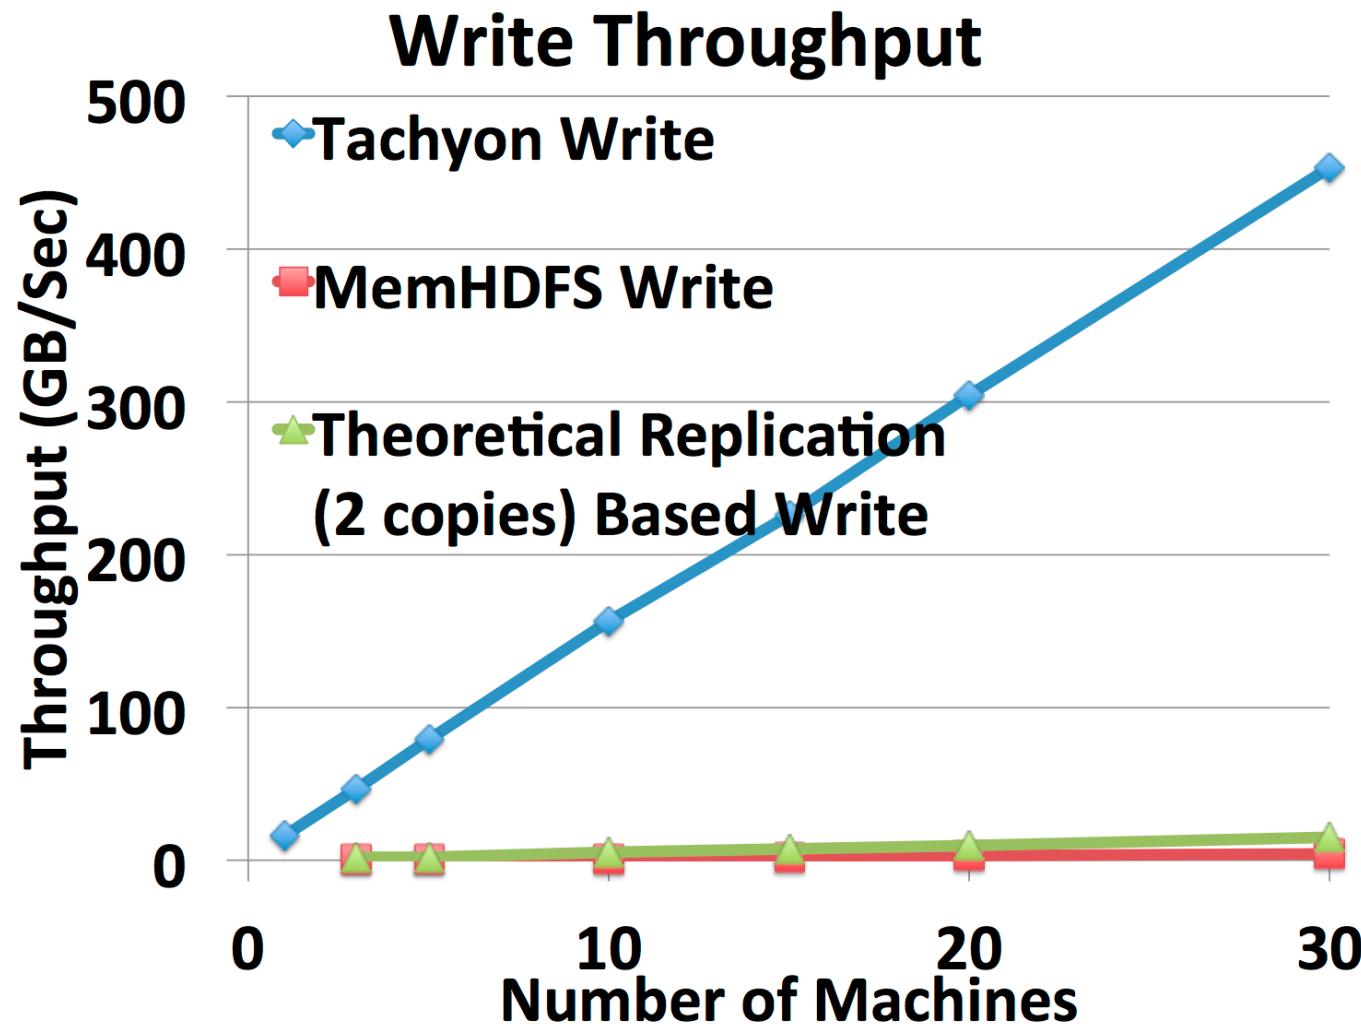
***No in-memory data duplication,
much less GC***



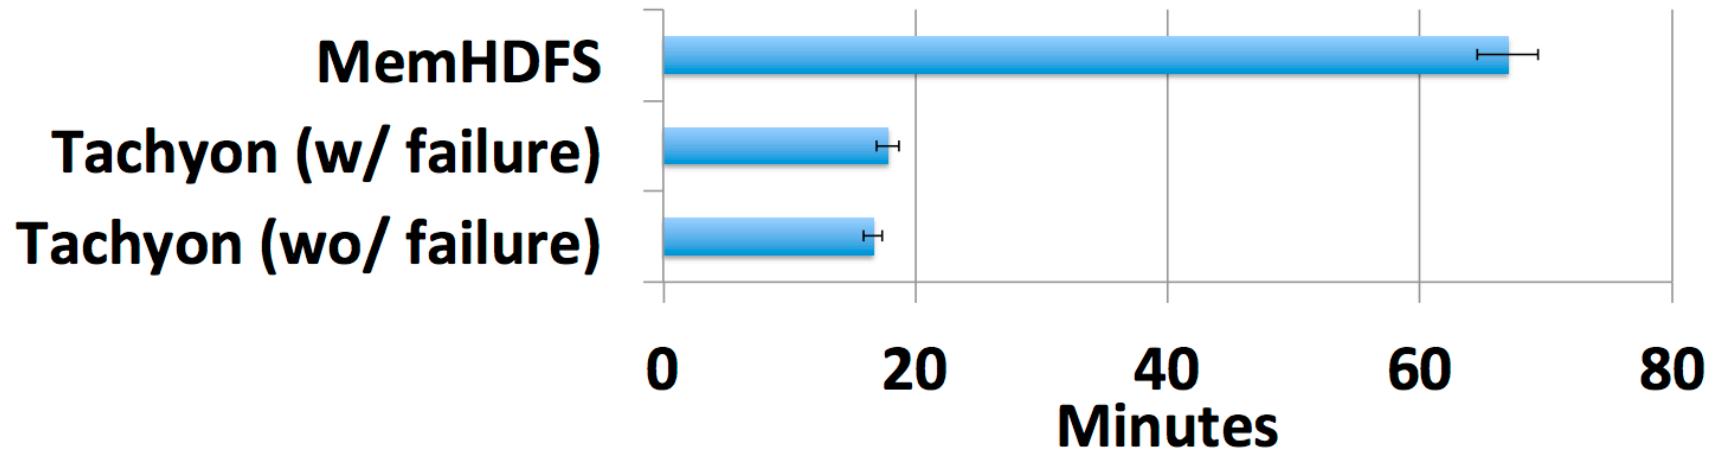
Lineage in Storage (*alpha*)



Comparison with in Memory HDFS



Workflow Improvement



Performance comparison for realistic workflow. The workflow ran 4x faster on Tachyon than on MemHDFS. In case of node failure, applications in Tachyon still finishes 3.8x faster.

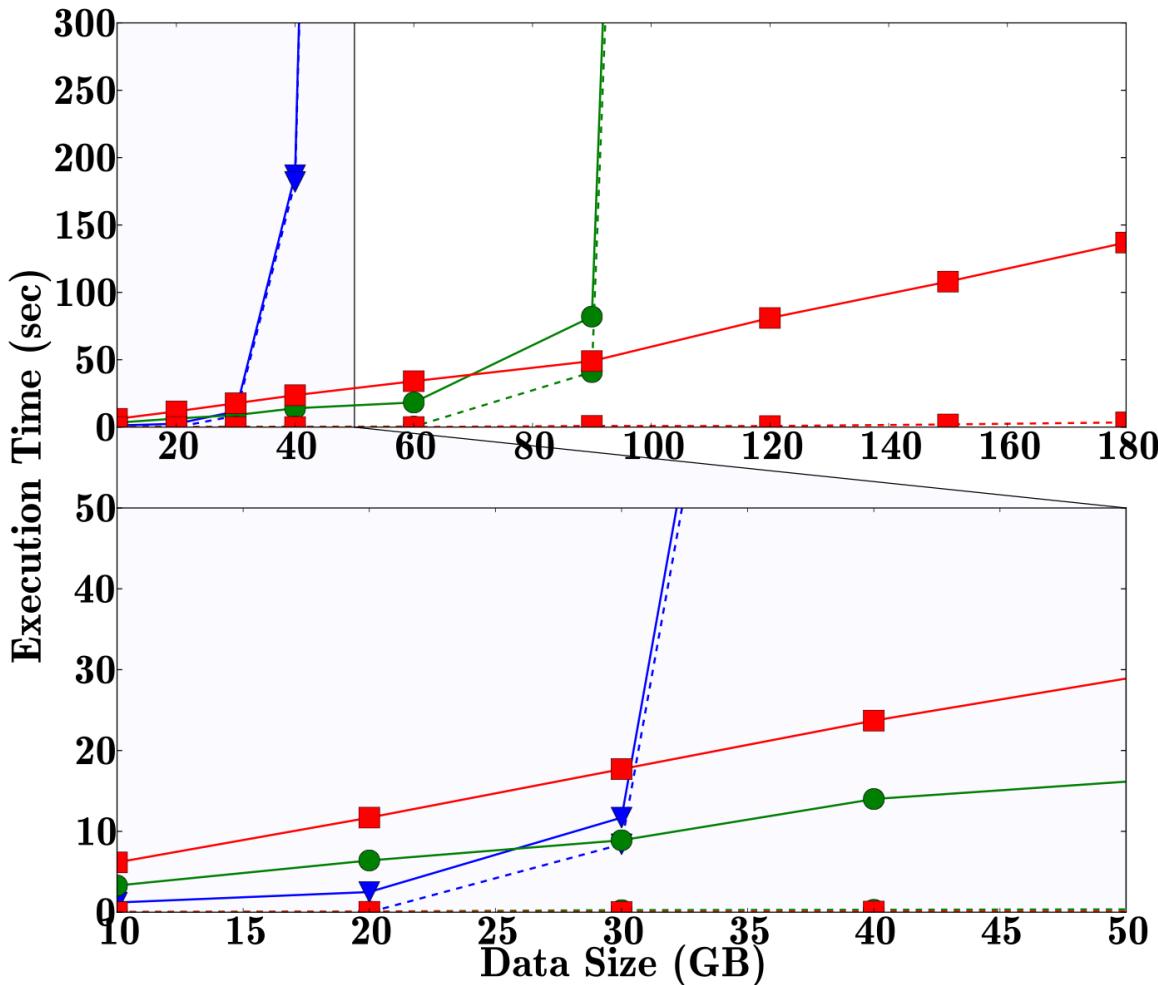
Further Improve Spark's Performance

▼—▼ Spark Cache (Deserialized)
▼--▼ Spark Cache GC (Deserialized)

●—● Spark Cache (Serialized)
●--● Spark Cache GC (Serialized)

■—■ TFS
■--■ TFS GC

Grep
Program



How easy / hard to use Tachyon?

Spark/MapReduce/Shark without Tachyon

- Spark
 - `val file = sc.textFile("hdfs://ip:port/path")`
- Hadoop MapReduce
 - `hadoop jar hadoop-examples-1.0.4.jar wordcount hdfs://localhost:19998/input hdfs://localhost:19998/output`
- Shark
 - `CREATE TABLE orders_cached AS SELECT * FROM orders;`

Spark/MapReduce/Shark with Tachyon

- Spark
 - val file = sc.textFile("tachyon://ip:port/path")
- Hadoop MapReduce
 - hadoop jar hadoop-examples-1.0.4.jar wordcount tachyon://localhost:19998/input tachyon://localhost:19998/output
- Shark
 - CREATE TABLE orders_tachyon AS SELECT * FROM orders;

Spark on Tachyon

```
./bin/spark-shell  
sc.hadoopConfiguration.set("fs.tachyon.impl", "tachyon.hadoop.TFS")
```

```
// Load input from Tachyon  
val file = sc.textFile("tachyon://localhost:19998/LICENSE")  
file.count() ; file.take(10);
```

```
// Store RDD OFF_HEAP in Tachyon  
import org.apache.spark.storage.StorageLevel;  
file.persist(StorageLevel.OFF_HEAP)  
file.count(); file.take(10);
```

```
// Save output to Tachyon  
file.flatMap(line => line.split(" ")).map(s => (s, 1)).reduceByKey((a, b) => a +  
b).saveAsTextFile("tachyon://localhost:19998/LICENSE_WC")
```

Outline

- Overview
 - Feature 1: Memory Centric Storage Architecture
 - Feature 2: Lineage in Storage
- Open Source
- Roadmap

History

Started at UC Berkeley AMPLab from the summer of 2012



- Reliable, Memory Speed Storage for Cluster Computing Frameworks
(UC Berkeley EECS Tech Report)
- Haoyuan Li, Ali Ghodsi, Matei Zaharia, Ion Stoica, Scott Shenker



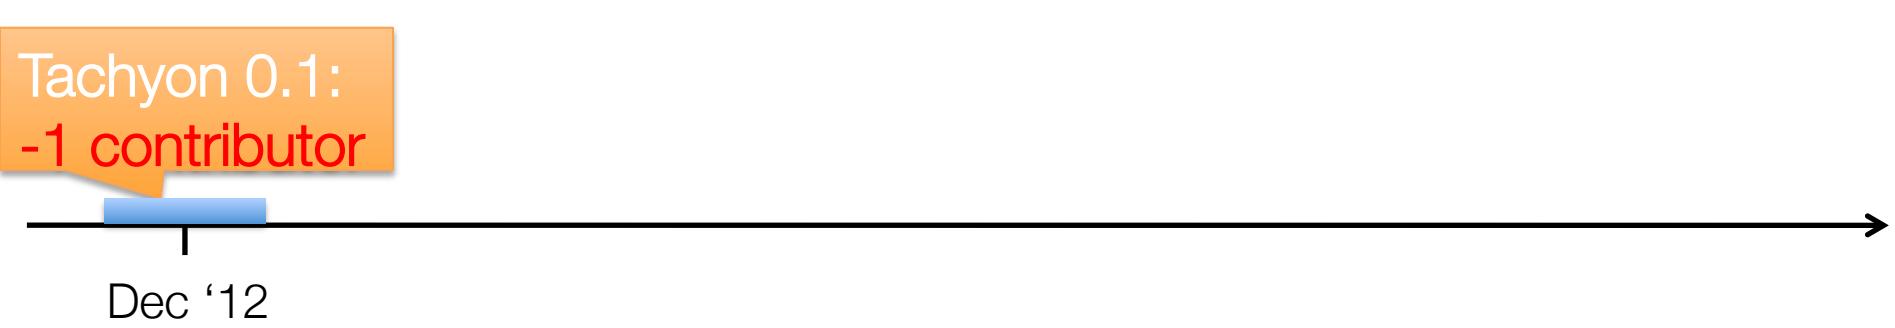
TACHYON Open Source Status

- Apache License 2.0, Version 0.5.0 (July 2014)



- Deployed at tens of companies
- 20+ Companies Contributing
- Spark and MapReduce applications can run without any code change

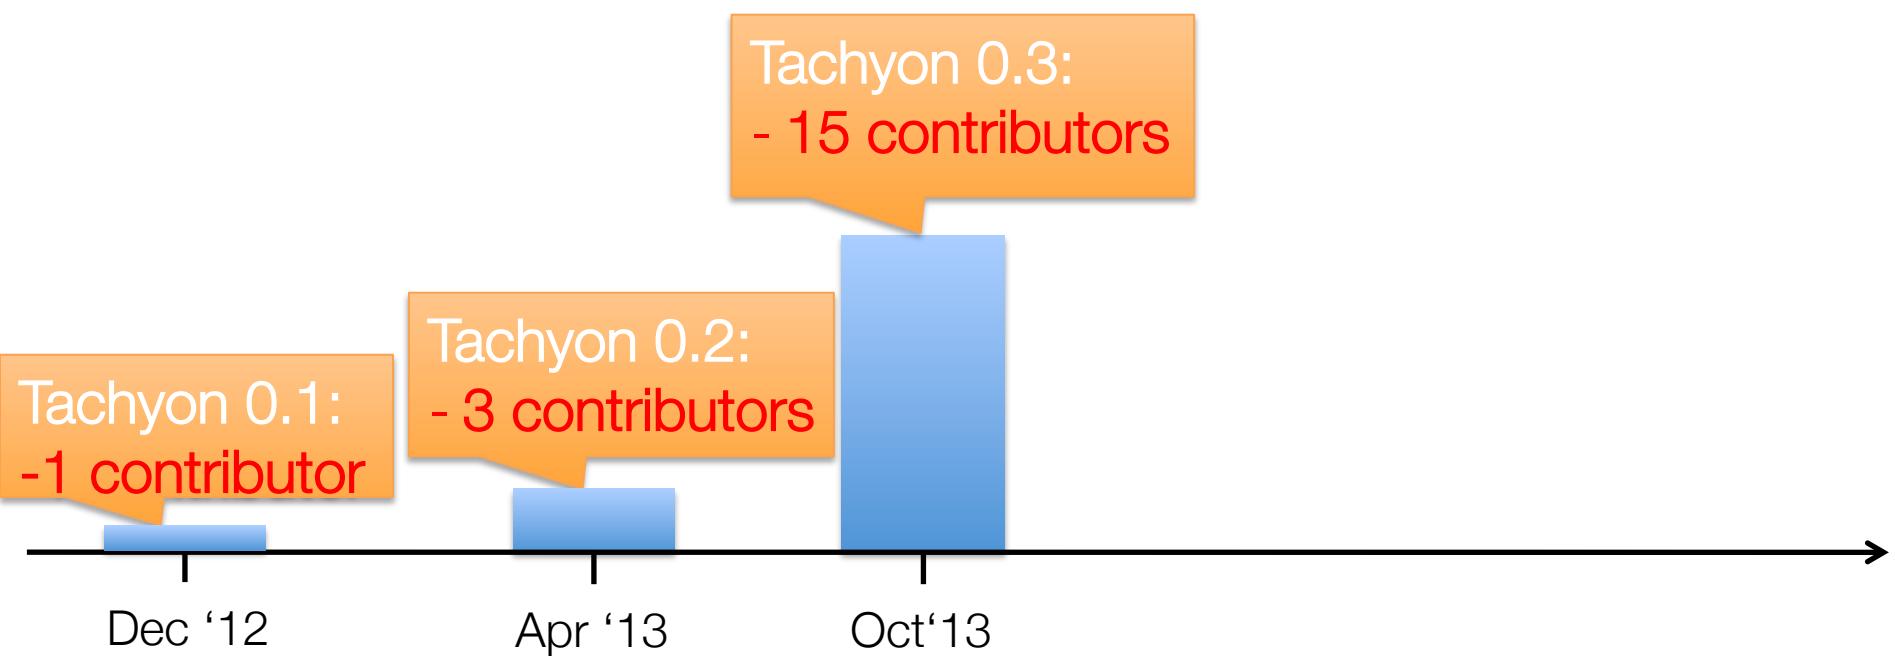
Release Growth



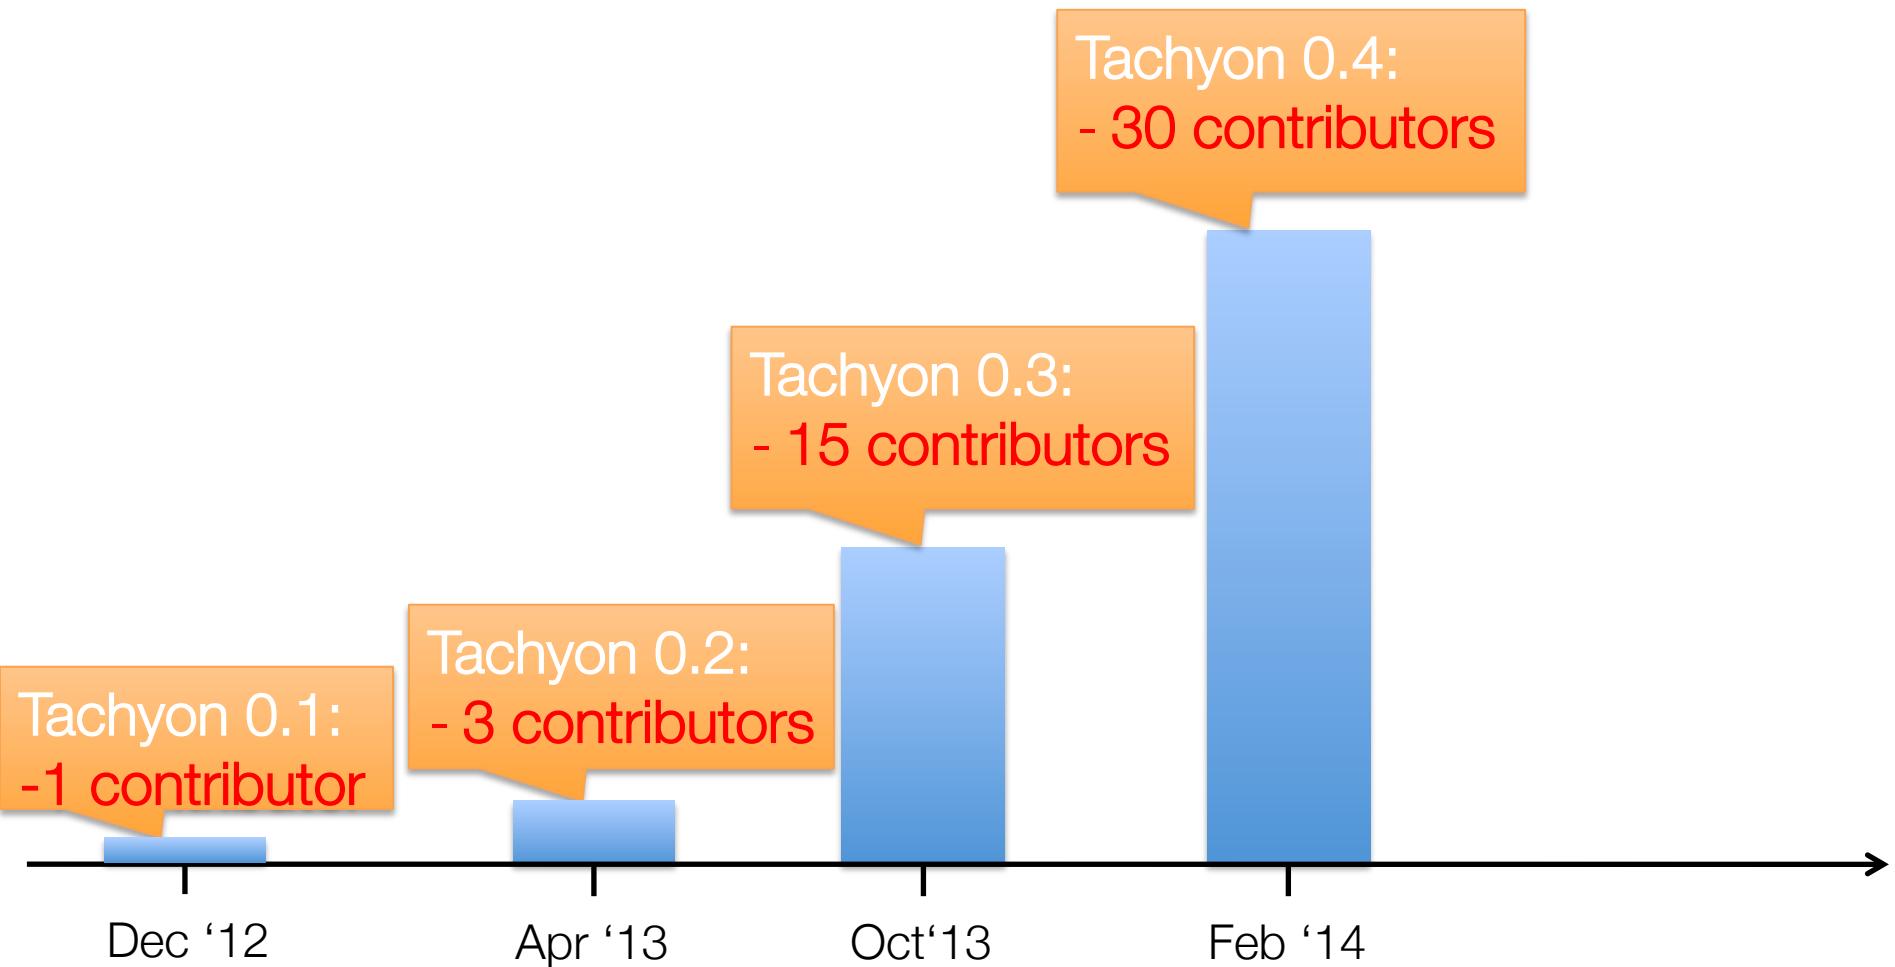
Release Growth



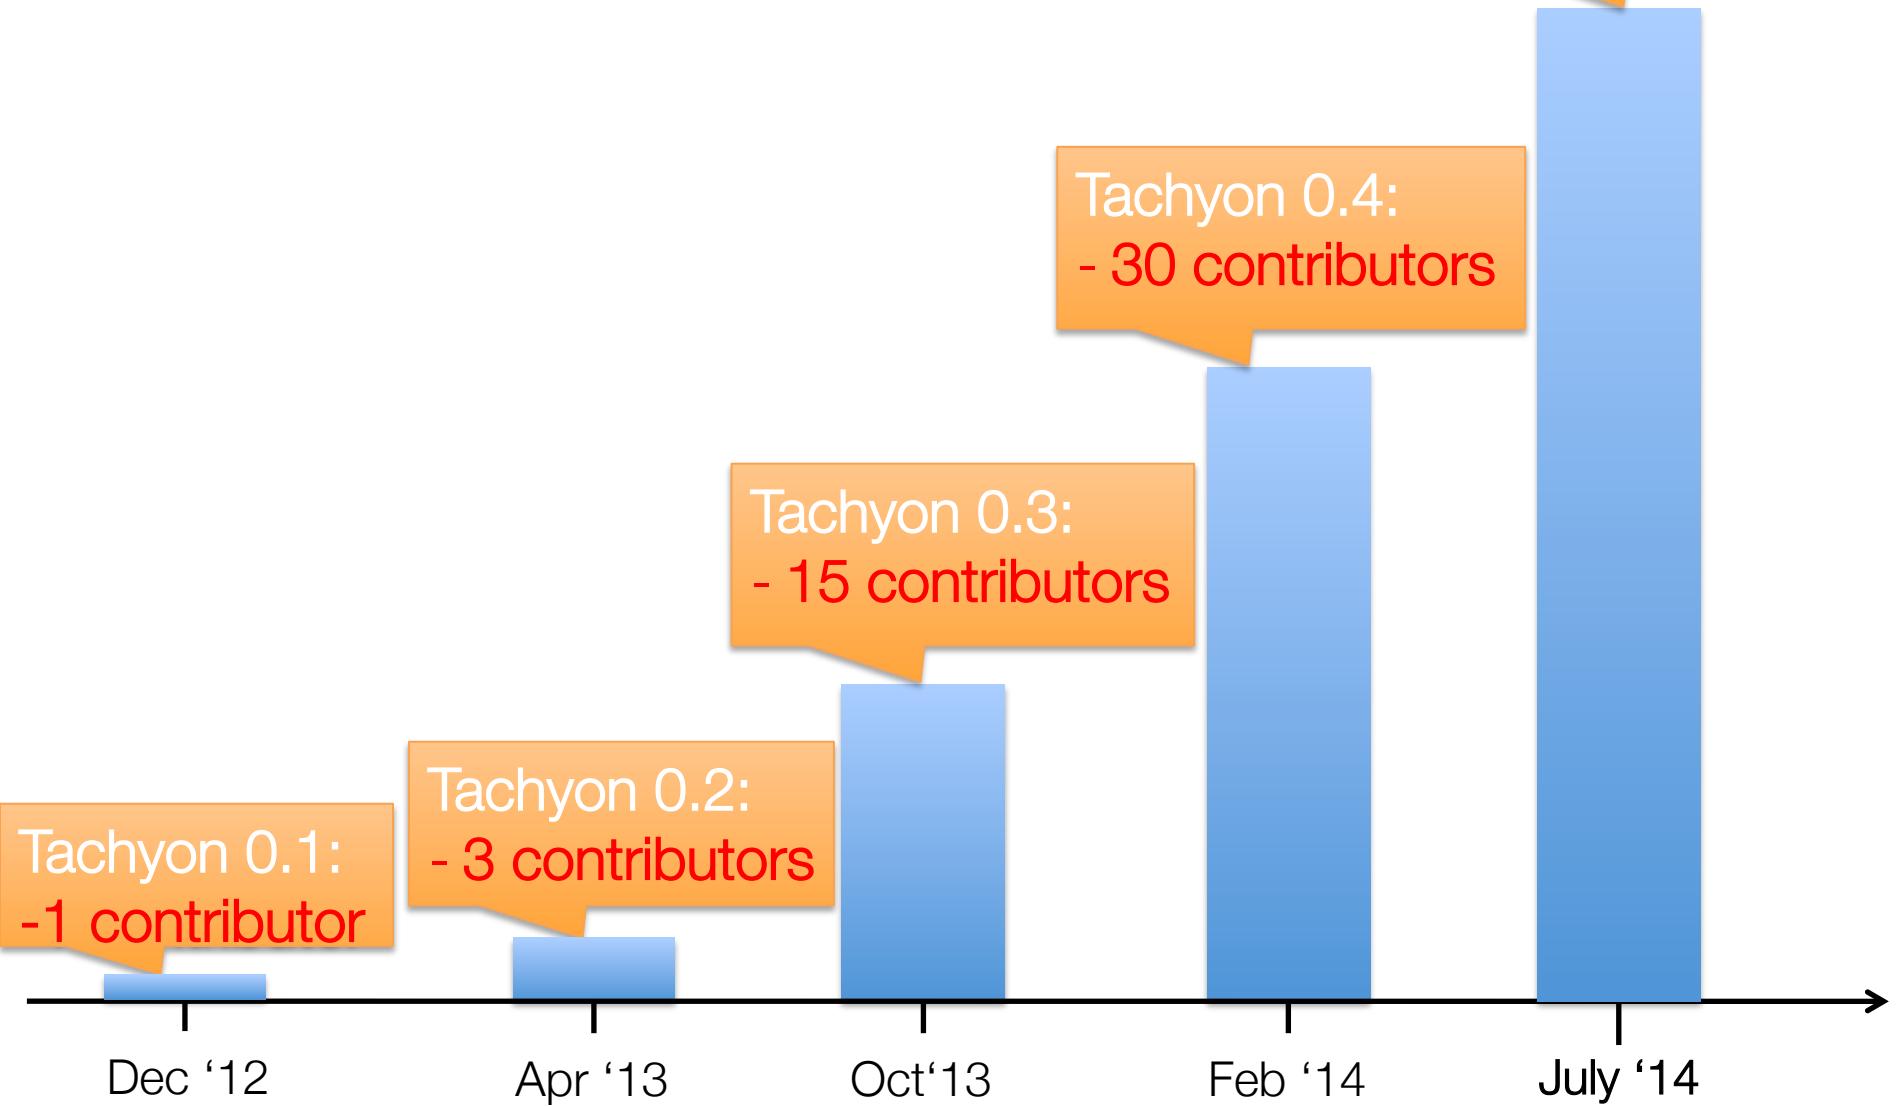
Release Growth



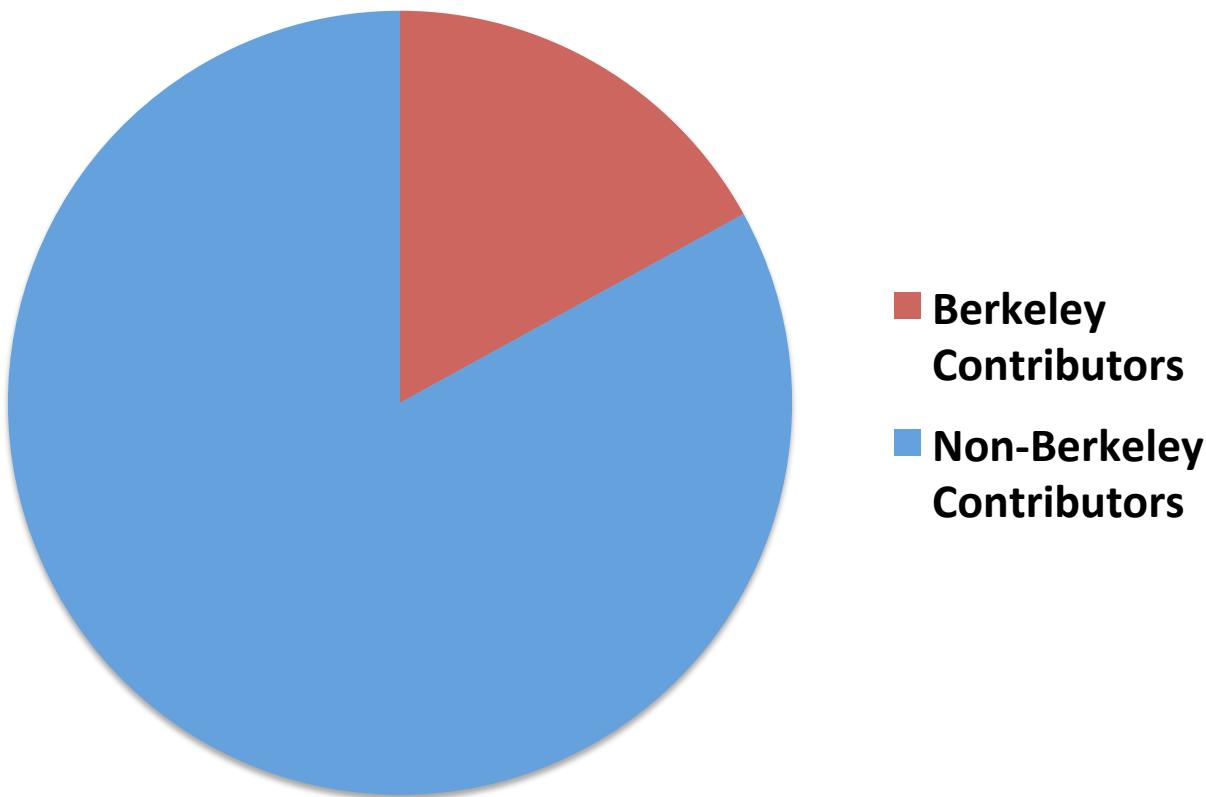
Release Growth



Release Growth



Open Community



Thanks to our Code Contributors!

Aaron Davidson

Achal Soni

Ali Ghodsi

Andrew Ash

Anurag Khandelwal

Aslan Bekirov

Bill Zhao

Brad Childs

Calvin Jia

Chao Chen

Cheng Chang

Cheng Hao

Colin Patrick McCabe

David Capwell

David Zhu

Du Li

Fei Wang

Gerald Zhang

Grace Huang

Haoyuan Li

Henry Saputra

Hobin Yoon

Huamin Chen

Jey Kottalam

Joseph Tang

Juan Zhou

Jun Aoki

Lin Xing

Lukasz Jastrzebski

Manu Goyal

Mark Hamstra

Mingfei Shi

Mubarak Seyed

Nick Lanham

Orcun Simsek

Pengfei Xuan

Qianhao Dong

Qifan Pu

Raymond Liu

Reynold Xin

Robert Metzger

Rong Gu

Sean Zhong

Seonghwan Moon

Shivaram Venkataraman

Srinivas Parayya

Tao Wang

Timothy St. Clair

Thu Kyaw

Vamsi Chittlers

Xi Liu

Xiang Zhong

Xiaomin Zhang

Zhao Zhang



Tachyon
is in Fedora 20

Thanks to Redhat!

Commercially supported

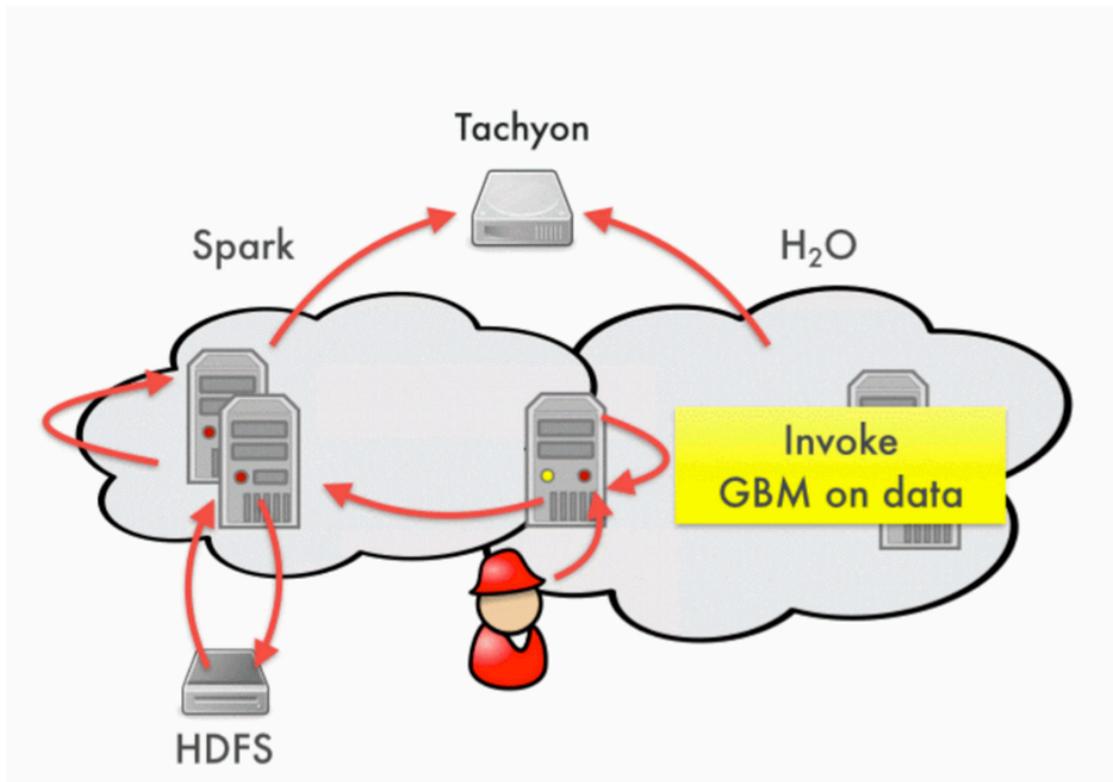


**and running in dozens of
their customers' clusters**

Tachyon is the
Default Off-Heap Storage
Solution for  Spark

Exchange Data Between Spark and H2O

Today, data gets parsed and exchanged between Spark and H2O via Tachyon. Users can interactively query big data both via SQL and ML from within the same context.



Believe from Industry

The Future Architecture of a Data Lake: In-memory Data Exchange Platform Using Tachyon and Apache Spark

OCTOBER 14, 2014 | NEWS | BY PAUL M. DAVIS



Pivotal and EMC are betting on
Spark cousin Tachyon as in-
memory file system

by [Derrick Harris](#) OCT. 14, 2014 - 11:47 AM PDT



Pivotal bets on Tachyon as next in- memory file system



Pivotal Expands on Data Lake Vision with Embrace of Project Tachyon

Oct 14, 2014

Reaching wider communities: e.g. GlusterFS

 blog.gluster.org/2014/08/glusterfs-and-tachyon/

 GLUSTER
COMMUNITY

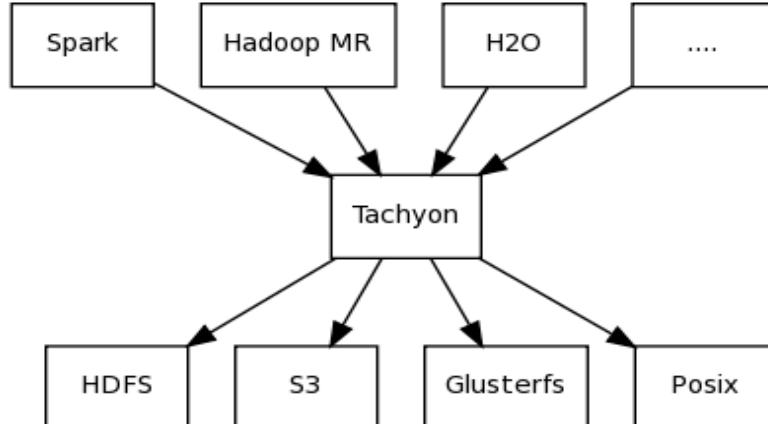
[DOCUMENTATION](#) [CONTACT](#) [ABOUT](#)

[Like](#) 1.1k [g+1](#) [Tweet](#)

by huamin on August 7, 2014 [← Return to Blog Home](#)

Glusterfs and Tachyon

Tachyon, an in-memory distributed filesystem, is among the most dynamic projects in big data analytics stack. It provides java io like API, support Apache Spark, and vastly improves Spark's performance under large data set.



```
graph TD; Spark[Spark] --> Tachyon[Tachyon]; HadoopMR[Hadoop MR] --> Tachyon; H2O[H2O] --> Tachyon; Ellipsis[...] --> Tachyon; Tachyon --> HDFS[HDFS]; Tachyon --> S3[S3]; Tachyon --> Glusterfs[Glusterfs]; Tachyon --> Posix[Posix]
```

The diagram illustrates the integration of Tachyon with various big data systems. At the top, four boxes represent different systems: Spark, Hadoop MR, H2O, and an ellipsis (...). Arrows point from each of these boxes down to a central box labeled "Tachyon". From the "Tachyon" box, arrows point down to four other boxes at the bottom: HDFS, S3, Glusterfs, and Posix.

Under Filesystem Choices (Big Data, Cloud, HPC, Enterprise)



Under Filesystem Choices (Big Data, Cloud, HPC, Enterprise)



Outline

- Overview
 - Feature 1: Memory Centric Storage Architecture
 - Feature 2: Lineage in Storage
- Open Source
- Roadmap

Features

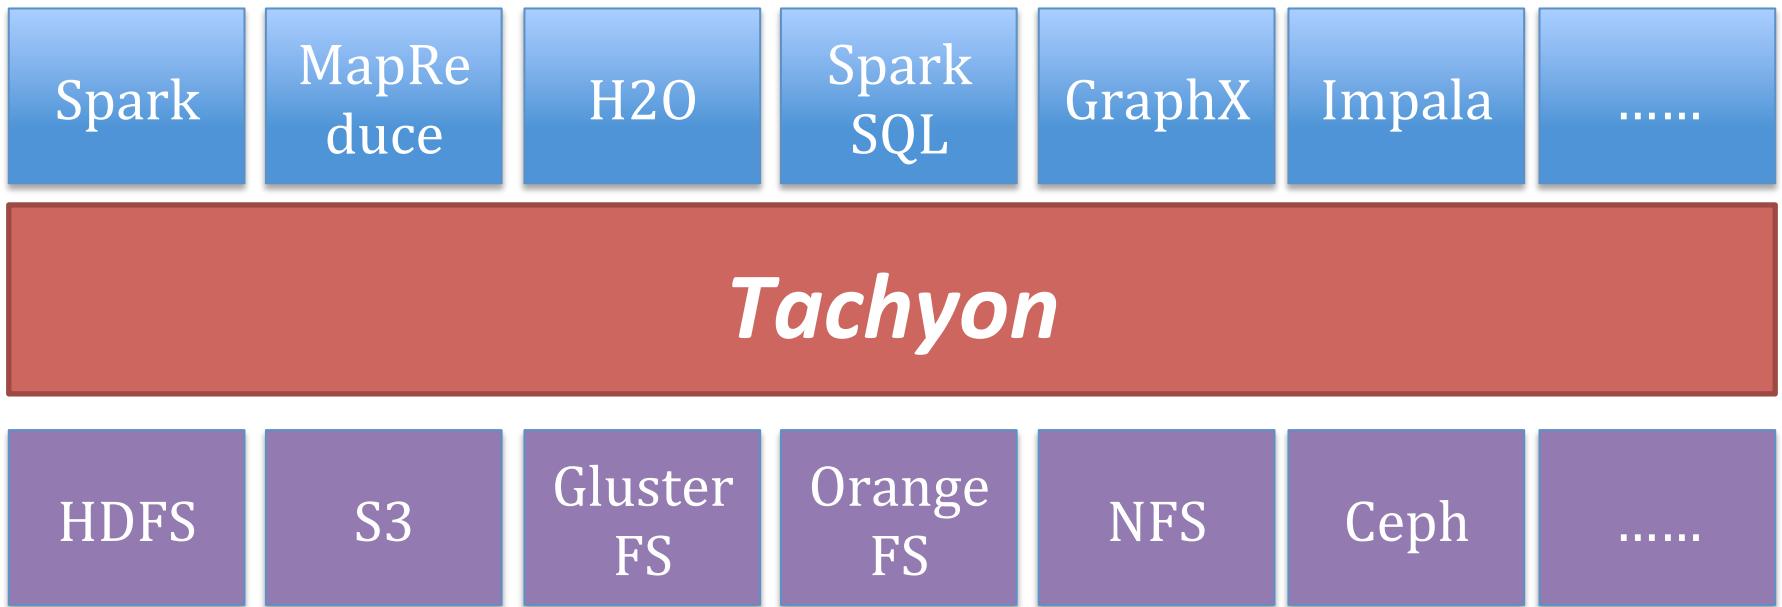
- Memory Centric Storage Architecture
- Lineage in Storage (alpha)
- Hierarchical Local Storage
- Data Serving
- Different hardware
- More...
- Your Requirements?

Short Term Roadmap (0.6 Release)

- Ceph Integration (Ceph Community, Redhat)
- Hierarchical Local Storage (Intel)
- Performance Improvement (Yahoo)
- Multi-tenancy (AMPLab)
- Mesos Integration (Mesos Community, Mesosphere)
- Network Sub-system Improvement (Pivotal)
- ***Many more*** from AMPLab and Industry Contributors

Goal?

Better Assist Other Components



Welcome Collaboration!

Thanks!

Questions?

- *More Information:*
 - Website: <http://tachyon-project.org>
 - Github: <https://github.com/amplab/tachyon>
 - Meetup: <http://www.meetup.com/Tachyon>
- *Email:* haoyuan@cs.berkeley.edu

Tachyon 0.5:
- 46 contributors

Release Growth

